

```
In [49]: import pandas as pd  
data=pd.read_csv("/home/placement/Downloads/fiat500 (another copy).csv")
```

```
In [50]: data.describe()
```

```
Out[50]:
```

|       | ID          | engine_power | age_in_days | km            | previous_owners | lat         | lon         | price        |
|-------|-------------|--------------|-------------|---------------|-----------------|-------------|-------------|--------------|
| count | 1538.000000 | 1538.000000  | 1538.000000 | 1538.000000   | 1538.000000     | 1538.000000 | 1538.000000 | 1538.000000  |
| mean  | 769.500000  | 51.904421    | 1650.980494 | 53396.011704  | 1.123537        | 43.541361   | 11.563428   | 8576.003901  |
| std   | 444.126671  | 3.988023     | 1289.522278 | 40046.830723  | 0.416423        | 2.133518    | 2.328190    | 1939.958641  |
| min   | 1.000000    | 51.000000    | 366.000000  | 1232.000000   | 1.000000        | 36.855839   | 7.245400    | 2500.000000  |
| 25%   | 385.250000  | 51.000000    | 670.000000  | 20006.250000  | 1.000000        | 41.802990   | 9.505090    | 7122.500000  |
| 50%   | 769.500000  | 51.000000    | 1035.000000 | 39031.000000  | 1.000000        | 44.394096   | 11.869260   | 9000.000000  |
| 75%   | 1153.750000 | 51.000000    | 2616.000000 | 79667.750000  | 1.000000        | 45.467960   | 12.769040   | 10000.000000 |
| max   | 1538.000000 | 77.000000    | 4658.000000 | 235000.000000 | 4.000000        | 46.795612   | 18.365520   | 11100.000000 |

```
In [51]: data2=data.drop(['ID','lat','lon'],axis=1)
```

In [52]: data2

Out[52]:

|      | model  | engine_power | age_in_days | km     | previous_owners | price |
|------|--------|--------------|-------------|--------|-----------------|-------|
| 0    | lounge | 51           | 882         | 25000  | 1               | 8900  |
| 1    | pop    | 51           | 1186        | 32500  | 1               | 8800  |
| 2    | sport  | 74           | 4658        | 142228 | 1               | 4200  |
| 3    | lounge | 51           | 2739        | 160000 | 1               | 6000  |
| 4    | pop    | 73           | 3074        | 106880 | 1               | 5700  |
| ...  | ...    | ...          | ...         | ...    | ...             | ...   |
| 1533 | sport  | 51           | 3712        | 115280 | 1               | 5200  |
| 1534 | lounge | 74           | 3835        | 112000 | 1               | 4600  |
| 1535 | pop    | 51           | 2223        | 60457  | 1               | 7500  |
| 1536 | lounge | 51           | 2557        | 80750  | 1               | 5990  |
| 1537 | pop    | 51           | 1766        | 54276  | 1               | 7900  |

1538 rows × 6 columns

In [53]: data2=pd.get\_dummies(data2)

In [54]: data2

Out[54]:

|      | engine_power | age_in_days | km     | previous_owners | price | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|-------|--------------|-----------|-------------|
| 0    | 51           | 882         | 25000  | 1               | 8900  | 1            | 0         | 0           |
| 1    | 51           | 1186        | 32500  | 1               | 8800  | 0            | 1         | 0           |
| 2    | 74           | 4658        | 142228 | 1               | 4200  | 0            | 0         | 1           |
| 3    | 51           | 2739        | 160000 | 1               | 6000  | 1            | 0         | 0           |
| 4    | 73           | 3074        | 106880 | 1               | 5700  | 0            | 1         | 0           |
| ...  | ...          | ...         | ...    | ...             | ...   | ...          | ...       | ...         |
| 1533 | 51           | 3712        | 115280 | 1               | 5200  | 0            | 0         | 1           |
| 1534 | 74           | 3835        | 112000 | 1               | 4600  | 1            | 0         | 0           |
| 1535 | 51           | 2223        | 60457  | 1               | 7500  | 0            | 1         | 0           |
| 1536 | 51           | 2557        | 80750  | 1               | 5990  | 1            | 0         | 0           |
| 1537 | 51           | 1766        | 54276  | 1               | 7900  | 0            | 1         | 0           |

1538 rows × 8 columns

In [55]: data2.shape

Out[55]: (1538, 8)

In [56]: data2

Out[56]:

|      | engine_power | age_in_days | km     | previous_owners | price | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|-------|--------------|-----------|-------------|
| 0    | 51           | 882         | 25000  | 1               | 8900  | 1            | 0         | 0           |
| 1    | 51           | 1186        | 32500  | 1               | 8800  | 0            | 1         | 0           |
| 2    | 74           | 4658        | 142228 | 1               | 4200  | 0            | 0         | 1           |
| 3    | 51           | 2739        | 160000 | 1               | 6000  | 1            | 0         | 0           |
| 4    | 73           | 3074        | 106880 | 1               | 5700  | 0            | 1         | 0           |
| ...  | ...          | ...         | ...    | ...             | ...   | ...          | ...       | ...         |
| 1533 | 51           | 3712        | 115280 | 1               | 5200  | 0            | 0         | 1           |
| 1534 | 74           | 3835        | 112000 | 1               | 4600  | 1            | 0         | 0           |
| 1535 | 51           | 2223        | 60457  | 1               | 7500  | 0            | 1         | 0           |
| 1536 | 51           | 2557        | 80750  | 1               | 5990  | 1            | 0         | 0           |
| 1537 | 51           | 1766        | 54276  | 1               | 7900  | 0            | 1         | 0           |

1538 rows × 8 columns

In [57]: `y=data2['price']`  
`x=data2.drop('price',axis=1)`

In [58]:

```
y
```

Out[58]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
Name: price, Length: 1538, dtype: int64
```

In [59]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.33,random_state=42)
```

In [60]:

```
x_test.head(5)
```

Out[60]:

|      | engine_power | age_in_days | km     | previous_owners | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|--------------|-----------|-------------|
| 481  | 51           | 3197        | 120000 | 2               | 0            | 1         | 0           |
| 76   | 62           | 2101        | 103000 | 1               | 0            | 1         | 0           |
| 1502 | 51           | 670         | 32473  | 1               | 1            | 0         | 0           |
| 669  | 51           | 913         | 29000  | 1               | 1            | 0         | 0           |
| 1409 | 51           | 762         | 18800  | 1               | 1            | 0         | 0           |

```
In [61]: x_train.head(5)
```

```
Out[61]:
```

|     | engine_power | age_in_days | km    | previous_owners | model_lounge | model_pop | model_sport |
|-----|--------------|-------------|-------|-----------------|--------------|-----------|-------------|
| 527 | 51           | 425         | 13111 | 1               | 1            | 0         | 0           |
| 129 | 51           | 1127        | 21400 | 1               | 1            | 0         | 0           |
| 602 | 51           | 2039        | 57039 | 1               | 0            | 1         | 0           |
| 331 | 51           | 1155        | 40700 | 1               | 1            | 0         | 0           |
| 323 | 51           | 425         | 16783 | 1               | 1            | 0         | 0           |

```
In [62]: y_test.head(5)
```

```
Out[62]: 481      7900
76       7900
1502     9400
669      8500
1409     9700
Name: price, dtype: int64
```

```
In [63]: y_train.head(5)
```

```
Out[63]: 527      9990
129      9500
602      7590
331      8750
323      9100
Name: price, dtype: int64
```

```
In [64]: x_train.shape
```

```
Out[64]: (1030, 7)
```

In [65]:

```
y_train
```

Out[65]:

|     |      |
|-----|------|
| 527 | 9990 |
| 129 | 9500 |
| 602 | 7590 |
| 331 | 8750 |
| 323 | 9100 |

|      |       |
|------|-------|
|      | ...   |
| 1130 | 10990 |
| 1294 | 9800  |
| 860  | 5500  |
| 1459 | 9990  |
| 1126 | 8900  |

Name: price, Length: 1030, dtype: int64

In [66]:

```
from sklearn.linear_model import LinearRegression  
reg=LinearRegression()  
reg.fit(x_train,y_train)
```

Out[66]:

```
LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [67]:

```
ypred=reg.predict(x_test)
```

In [68]:

ypred

Out[68]: array([ 5867.6503378 , 7133.70142341, 9866.35776216, 9723.28874535,  
10039.59101162, 9654.07582608, 9673.14563045, 10118.70728123,  
9903.85952664, 9351.55828437, 10434.34963575, 7732.26255693,  
7698.67240131, 6565.95240435, 9662.90103518, 10373.20344286,  
9599.94844451, 7699.34400418, 4941.33017994, 10455.2719478 ,  
10370.51555682, 10391.60424404, 7529.06622456, 9952.37340054,  
7006.13845729, 9000.1780961 , 4798.36770637, 6953.10376491,  
7810.39767825, 9623.80497535, 7333.52158317, 5229.18705519,  
5398.21541073, 5157.65652129, 8948.63632836, 5666.62365159,  
9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,  
9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,  
8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,  
7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,  
10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,  
9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,  
10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,  
6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,  
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,  
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,  
8828.78555188, 10025.02571528, 7278.77100022, 8411.45004006])

In [73]: `from sklearn.metrics import r2_score`  
`r2_score(y_test,ypred)`

Out[73]: 0.8415526986865394

In [74]: `from sklearn.metrics import mean_squared_error`  
`mean_squared_error(ypred,y_test)`

Out[74]: 581887.727391353

In [79]: `import math`  
`y=math.sqrt(581887)`

In [80]: `y`

Out[80]: 762.815180761369



```
In [81]: (581887)**(1/2)
```

```
Out[81]: 762.815180761369
```

```
In [85]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

```
Out[85]:
```

|    | index | price | predicted    | ID |
|----|-------|-------|--------------|----|
| 0  | 481   | 7900  | 5867.650338  | 0  |
| 1  | 76    | 7900  | 7133.701423  | 1  |
| 2  | 1502  | 9400  | 9866.357762  | 2  |
| 3  | 669   | 8500  | 9723.288745  | 3  |
| 4  | 1409  | 9700  | 10039.591012 | 4  |
| 5  | 1414  | 9900  | 9654.075826  | 5  |
| 6  | 1089  | 9900  | 9673.145630  | 6  |
| 7  | 1507  | 9950  | 10118.707281 | 7  |
| 8  | 970   | 10700 | 9903.859527  | 8  |
| 9  | 1198  | 8999  | 9351.558284  | 9  |
| 10 | 1088  | 9890  | 10434.349636 | 10 |
| 11 | 576   | 7990  | 7732.262557  | 11 |
| 12 | 965   | 7380  | 7698.672401  | 12 |
| 13 | 1488  | 6800  | 6565.952404  | 13 |
| 14 | 1432  | 8900  | 9662.901035  | 14 |

```
In [86]: Results['diff']=Results.apply(lambda row: row.price-row.predicted,axis=1)
```

In [87]:

Results

Out[87]:

|            | index | price | predicted    | ID  | diff         |
|------------|-------|-------|--------------|-----|--------------|
| <b>0</b>   | 481   | 7900  | 5867.650338  | 0   | 2032.349662  |
| <b>1</b>   | 76    | 7900  | 7133.701423  | 1   | 766.298577   |
| <b>2</b>   | 1502  | 9400  | 9866.357762  | 2   | -466.357762  |
| <b>3</b>   | 669   | 8500  | 9723.288745  | 3   | -1223.288745 |
| <b>4</b>   | 1409  | 9700  | 10039.591012 | 4   | -339.591012  |
| ...        | ...   | ...   | ...          | ... | ...          |
| <b>503</b> | 291   | 10900 | 10032.665135 | 503 | 867.334865   |
| <b>504</b> | 596   | 5699  | 6281.536277  | 504 | -582.536277  |
| <b>505</b> | 1489  | 9500  | 9986.327508  | 505 | -486.327508  |
| <b>506</b> | 1436  | 6990  | 8381.517020  | 506 | -1391.517020 |
| <b>507</b> | 575   | 10900 | 10371.142553 | 507 | 528.857447   |

508 rows × 5 columns

In [ ]: